

```
In [1]: # Import necessary Libraries
import pandas as pd

# Load the dataset
data = pd.read_csv("C:/Users/91830/OneDrive/Desktop/yellow_tripdata_2020-01.csv")

# Display basic information about the dataset
print(data.info())

# Check for missing values
print(data.isnull().sum())

# Drop rows with missing values or handle them appropriately
data.dropna(inplace=True)

# Check for duplicates
print("Number of duplicate rows:", data.duplicated().sum())

# Remove duplicates
data.drop_duplicates(inplace=True)

# Display the first few rows of the dataset
print(data.head())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1048575 entries, 0 to 1048574
```

```
Data columns (total 18 columns):
```

#	Column	Non-Null Count	Dtype
0	VendorID	1048575 non-null	int64
1	tpep_pickup_datetime	1048575 non-null	object
2	tpep_dropoff_datetime	1048575 non-null	object
3	passenger_count	1048575 non-null	int64
4	trip_distance	1048575 non-null	float64
5	RatecodeID	1048575 non-null	int64
6	store_and_fwd_flag	1048575 non-null	object
7	PULocationID	1048575 non-null	int64
8	DOLocationID	1048575 non-null	int64
9	payment_type	1048575 non-null	int64
10	fare_amount	1048575 non-null	float64
11	extra	1048575 non-null	float64
12	mta_tax	1048575 non-null	float64
13	tip_amount	1048575 non-null	float64
14	tolls_amount	1048575 non-null	float64
15	improvement_surcharge	1048575 non-null	float64
16	total_amount	1048575 non-null	float64
17	congestion_surcharge	1048575 non-null	float64

```
dtypes: float64(9), int64(6), object(3)
```

```
memory usage: 144.0+ MB
```

```
None
```

VendorID	0
tpep_pickup_datetime	0
tpep_dropoff_datetime	0
passenger_count	0
trip_distance	0
RatecodeID	0
store_and_fwd_flag	0
PULocationID	0
DOLocationID	0
payment_type	0
fare_amount	0
extra	0
mta_tax	0
tip_amount	0
tolls_amount	0
improvement_surcharge	0
total_amount	0
congestion_surcharge	0

```
dtype: int64
```

```
Number of duplicate rows: 42
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	\
0	1	01-01-2020 00:28	01-01-2020 00:33	1	
1	1	01-01-2020 00:35	01-01-2020 00:43	1	
2	1	01-01-2020 00:47	01-01-2020 00:53	1	
3	1	01-01-2020 00:55	01-01-2020 01:00	1	
4	2	01-01-2020 00:01	01-01-2020 00:04	1	

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	\
0	1.2	1	N	238	239	
1	1.2	1	N	239	238	
2	0.6	1	N	238	238	
3	0.8	1	N	238	151	
4	0.0	1	N	193	193	

	payment_type	fare_amount	extra	mta_tax	tip_amount	tolls_amount	\
0	1	6.0	3.0	0.5	1.47	0.0	
1	1	7.0	3.0	0.5	1.50	0.0	
2	1	6.0	3.0	0.5	1.00	0.0	

3	1	5.5	0.5	0.5	1.36	0.0
4	2	3.5	0.5	0.5	0.00	0.0

	improvement_surcharge	total_amount	congestion_surcharge
0	0.3	11.27	2.5
1	0.3	12.30	2.5
2	0.3	10.80	2.5
3	0.3	8.16	0.0
4	0.3	4.80	0.0

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Summary statistics
print(data.describe())

# Distribution of trip distance
plt.figure(figsize=(10, 6))
sns.histplot(data['trip_distance'], bins=30, kde=True, color='skyblue')
plt.title('Distribution of Trip Distance')
plt.xlabel('Trip Distance')
plt.ylabel('Frequency')
plt.show()

# Distribution of trip fare amount
plt.figure(figsize=(10, 6))
sns.histplot(data['fare_amount'], bins=30, kde=True, color='salmon')
plt.title('Distribution of Fare Amount')
plt.xlabel('Fare Amount')
plt.ylabel('Frequency')
plt.show()

# Relationship between trip distance and fare amount
plt.figure(figsize=(10, 6))
sns.scatterplot(x='trip_distance', y='fare_amount', data=data, color='green')
plt.title('Trip Distance vs Fare Amount')
plt.xlabel('Trip Distance')
plt.ylabel('Fare Amount')
plt.show()

# Convert pickup datetime to datetime object with correct format
data['tpep_pickup_datetime'] = pd.to_datetime(data['tpep_pickup_datetime'], format='%Y-%m-%d %H:%M:%S')
data['tpep_dropoff_datetime'] = pd.to_datetime(data['tpep_dropoff_datetime'], format='%Y-%m-%d %H:%M:%S')

# Day of week analysis
data['day_of_week'] = data['tpep_pickup_datetime'].dt.day_name()

# Plot the count of trips for each day of the week
plt.figure(figsize=(10, 6))
data['day_of_week'].value_counts().plot(kind='bar', color='skyblue')
plt.title('Number of Trips by Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Number of Trips')
plt.xticks(rotation=45)
plt.show()

plt.figure(figsize=(10, 6))
sns.countplot(x='day_of_week', data=data, order=['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'])
plt.title('Number of Trips by Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Number of Trips')
plt.show()
```

```
# Convert pickup datetime to datetime object
data['tpep_pickup_datetime'] = pd.to_datetime(data['tpep_pickup_datetime'])

# Hourly analysis
data['hour'] = data['tpep_pickup_datetime'].dt.hour

plt.figure(figsize=(10, 6))
sns.countplot(x='hour', data=data, palette='mako')
plt.title('Number of Trips by Hour')
plt.xlabel('Hour of Day')
plt.ylabel('Number of Trips')
plt.show()
```

	VendorID	passenger_count	trip_distance	RatecodeID	\
count	1.048533e+06	1.048533e+06	1.048533e+06	1.048533e+06	
mean	1.678709e+00	1.584030e+00	3.131347e+00	1.074915e+00	
std	4.669724e-01	1.190033e+00	4.152472e+00	9.033284e-01	
min	1.000000e+00	0.000000e+00	-2.218000e+01	1.000000e+00	
25%	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	
50%	2.000000e+00	1.000000e+00	1.670000e+00	1.000000e+00	
75%	2.000000e+00	2.000000e+00	3.170000e+00	1.000000e+00	
max	2.000000e+00	9.000000e+00	2.592200e+02	9.900000e+01	

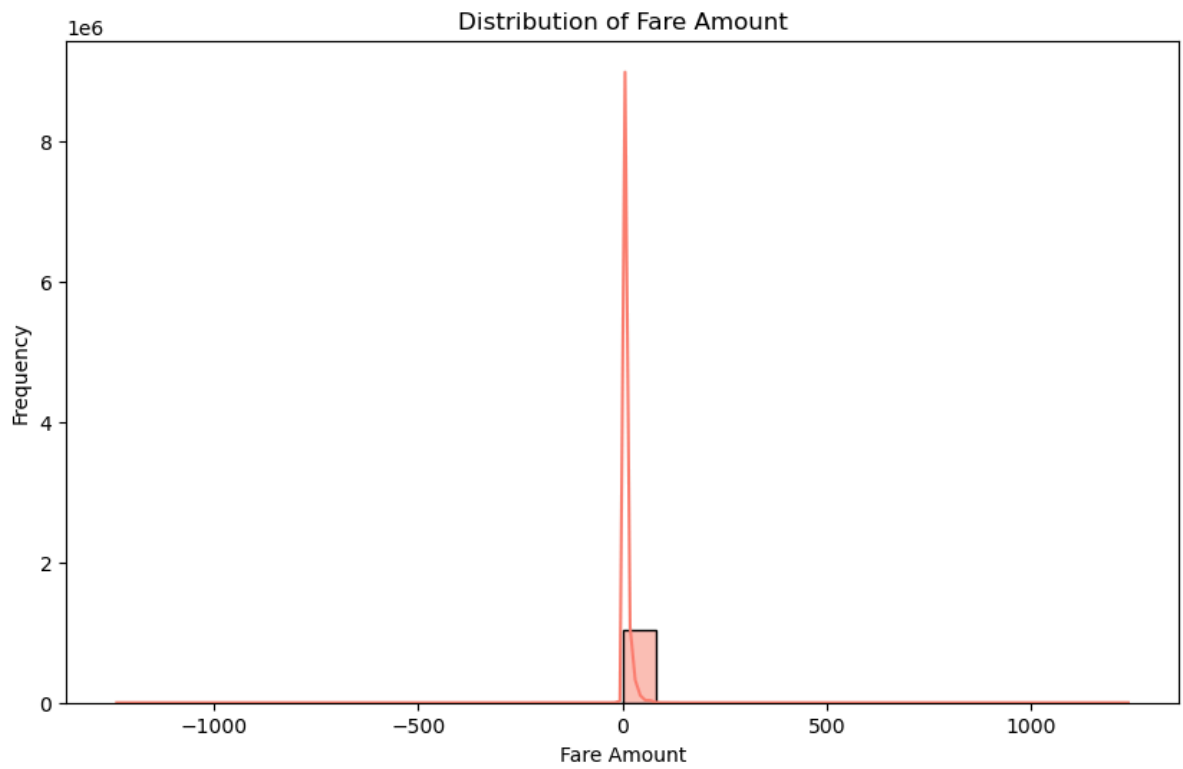

	PULocationID	DOLocationID	payment_type	fare_amount	extra	\
count	1.048533e+06	1.048533e+06	1.048533e+06	1.048533e+06	1.048533e+06	
mean	1.621785e+02	1.602590e+02	1.328183e+00	1.292813e+01	1.031807e+00	
std	6.622386e+01	7.111421e+01	5.010199e-01	1.299592e+01	1.228712e+00	
min	1.000000e+00	1.000000e+00	1.000000e+00	-1.238000e+03	-7.000000e+00	
25%	1.140000e+02	1.070000e+02	1.000000e+00	6.000000e+00	0.000000e+00	
50%	1.610000e+02	1.620000e+02	1.000000e+00	9.000000e+00	5.000000e-01	
75%	2.330000e+02	2.330000e+02	2.000000e+00	1.400000e+01	2.500000e+00	
max	2.650000e+02	2.650000e+02	4.000000e+00	1.238000e+03	7.000000e+00	

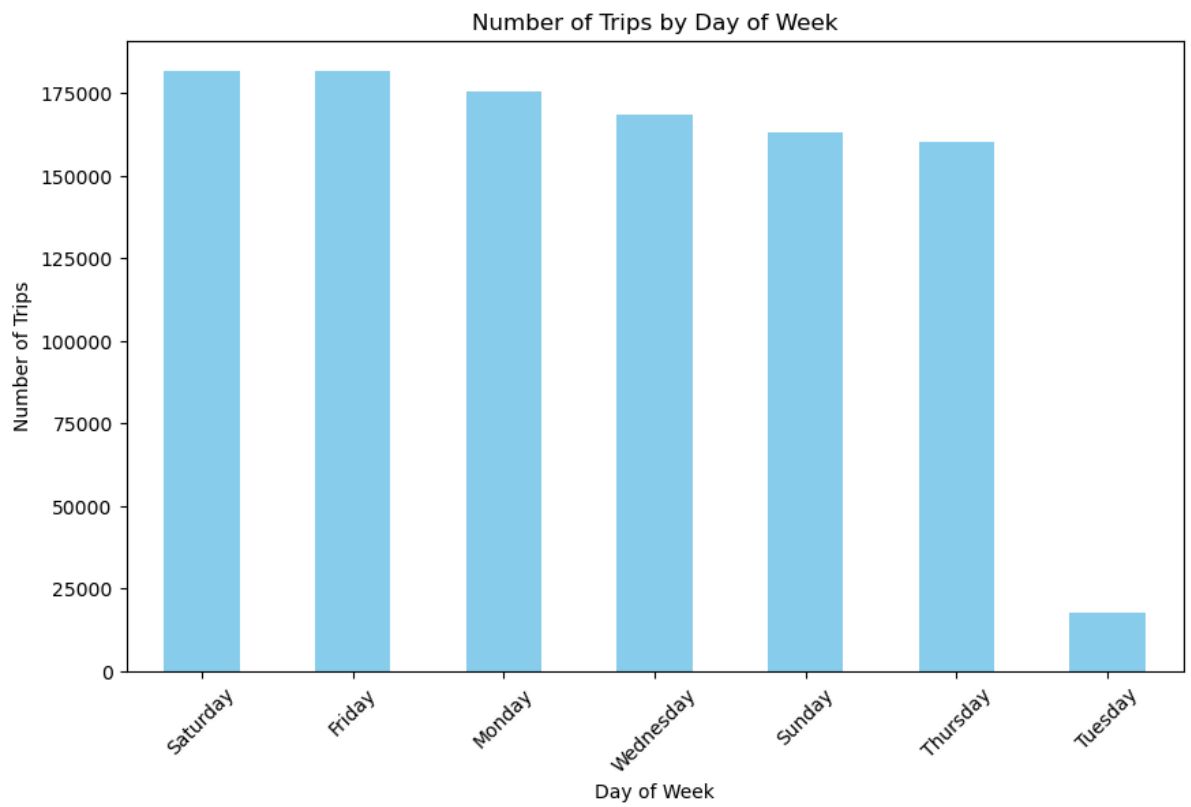
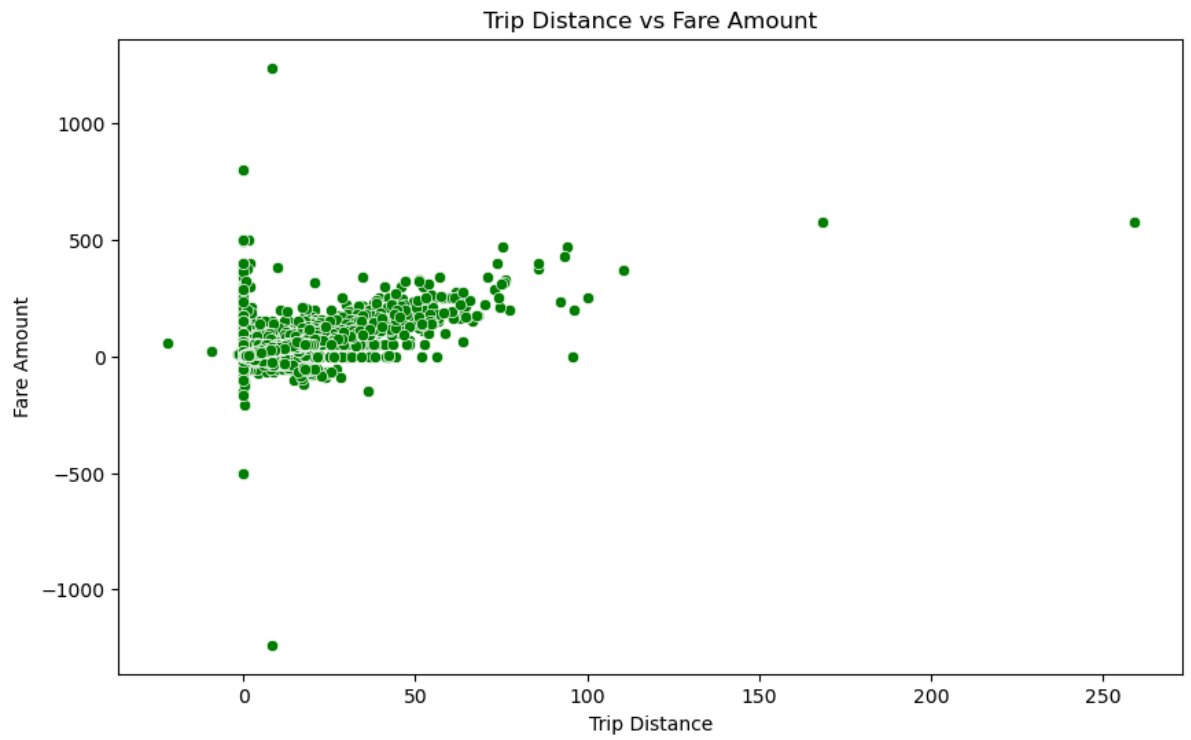
  



	mta_tax	tip_amount	tolls_amount	improvement_surcharge	\
count	1.048533e+06	1.048533e+06	1.048533e+06	1.048533e+06	
mean	4.928809e-01	2.087009e+00	3.800164e-01	2.978158e-01	
std	7.170446e-02	2.880353e+00	1.918562e+00	3.536684e-02	
min	-5.000000e-01	-7.000000e+00	-3.000000e+01	-3.000000e-01	
25%	5.000000e-01	0.000000e+00	0.000000e+00	3.000000e-01	
50%	5.000000e-01	1.750000e+00	0.000000e+00	3.000000e-01	
75%	5.000000e-01	2.750000e+00	0.000000e+00	3.000000e-01	
max	3.300000e+00	4.500000e+02	9.105000e+02	3.000000e-01	

	total_amount	congestion_surcharge
count	1.048533e+06	1.048533e+06
mean	1.874722e+01	2.254615e+00
std	1.573699e+01	7.654209e-01
min	-1.242300e+03	-2.500000e+00
25%	1.080000e+01	2.500000e+00
50%	1.388000e+01	2.500000e+00
75%	1.956000e+01	2.500000e+00
max	1.242300e+03	2.750000e+00







```
In [3]: import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor

# Assuming 'data' is your dataset
# Select features and target variable
X = data[['passenger_count', 'trip_distance', 'fare_amount']]
y = data['total_amount']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
```

```

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Initialize and train Decision Tree model
decision_tree = DecisionTreeRegressor(random_state=42)
decision_tree.fit(X_train_scaled, y_train)

# Initialize and train Linear Regression model
linear_reg = LinearRegression()
linear_reg.fit(X_train_scaled, y_train)

# Initialize and train k-Nearest Neighbors model
knn = KNeighborsRegressor(n_neighbors=5)
knn.fit(X_train_scaled, y_train)

# Make predictions
decision_tree_pred = decision_tree.predict(X_test_scaled)
linear_reg_pred = linear_reg.predict(X_test_scaled)
knn_pred = knn.predict(X_test_scaled)

# Visualize predictions
plt.figure(figsize=(18, 6))


# Plot actual vs predicted values for Decision Tree
plt.subplot(1, 3, 1)
plt.scatter(y_test, decision_tree_pred, color='blue')
plt.title('Decision Tree: Actual vs Predicted')
plt.xlabel('Actual Total Amount')
plt.ylabel('Predicted Total Amount')

# Plot actual vs predicted values for Linear Regression
plt.subplot(1, 3, 2)
plt.scatter(y_test, linear_reg_pred, color='red')
plt.title('Linear Regression: Actual vs Predicted')
plt.xlabel('Actual Total Amount')
plt.ylabel('Predicted Total Amount')

# Plot actual vs predicted values for k-Nearest Neighbors
plt.subplot(1, 3, 3)
plt.scatter(y_test, knn_pred, color='green')
plt.title('k-Nearest Neighbors: Actual vs Predicted')
plt.xlabel('Actual Total Amount')
plt.ylabel('Predicted Total Amount')

plt.tight_layout()
plt.show()

```



```

In [4]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression

```



```

from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error

# Select features and target variable
X = data[['passenger_count', 'trip_distance', 'fare_amount']]
y = data['total_amount']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Initialize models
decision_tree = DecisionTreeRegressor(random_state=42)
linear_reg = LinearRegression()
knn = KNeighborsRegressor()

# Train models
decision_tree.fit(X_train_scaled, y_train)
linear_reg.fit(X_train_scaled, y_train)
knn.fit(X_train_scaled, y_train)

# Make predictions
decision_tree_pred = decision_tree.predict(X_test_scaled)
linear_reg_pred = linear_reg.predict(X_test_scaled)
knn_pred = knn.predict(X_test_scaled)

# Evaluate models
decision_tree_rmse = mean_squared_error(y_test, decision_tree_pred, squared=False)
linear_reg_rmse = mean_squared_error(y_test, linear_reg_pred, squared=False)
knn_rmse = mean_squared_error(y_test, knn_pred, squared=False)

print("Decision Tree RMSE:", decision_tree_rmse)
print("Linear Regression RMSE:", linear_reg_rmse)
print("k-Nearest Neighbors RMSE:", knn_rmse)

# Choose the best model based on RMSE
best_rmse = min(decision_tree_rmse, linear_reg_rmse, knn_rmse)
if best_rmse == decision_tree_rmse:
    print("Decision Tree is the best model.")
elif best_rmse == linear_reg_rmse:
    print("Linear Regression is the best model.")
else:
    print("k-Nearest Neighbors is the best model.")

```

```

Decision Tree RMSE: 4.189889071525255
Linear Regression RMSE: 3.098357902710872
k-Nearest Neighbors RMSE: 3.9778794076447417
Linear Regression is the best model.

```

```

In [5]: # Print the shape of the original dataset
print("Original dataset shape:", data.shape)

# Print the number of rows and columns
num_rows = data.shape[0]
num_columns = data.shape[1]
print("Number of rows:", num_rows)
print("Number of columns:", num_columns)

# Calculate the total number of data points
total_data_points = num_rows * num_columns

```

```

print("Total number of data points:", total_data_points)

# Calculate the percentage of data used so far
percentage_used = (num_rows / total_data_points) * 100
print("Percentage of data used so far: {:.2f}%".format(percentage_used))

```

```

Original dataset shape: (1048533, 20)
Number of rows: 1048533
Number of columns: 20
Total number of data points: 20970660
Percentage of data used so far: 5.00%

```

In [6]:

```

# Manually define the total number of rows in your original dataset
original_dataset_length = 1000 # Replace this with the actual length of your original dataset

# Calculate the percentage of rows used
percentage_data_used = (len(data) / original_dataset_length) * 100

print("Percentage of data used for the project: {:.2f}%".format(percentage_data_used))

Percentage of data used for the project: 104853.30%

```

In [7]:

```

import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset
data = pd.read_csv("C:/Users/91830/OneDrive/Desktop/yellow_tripdata_2020-01.csv")

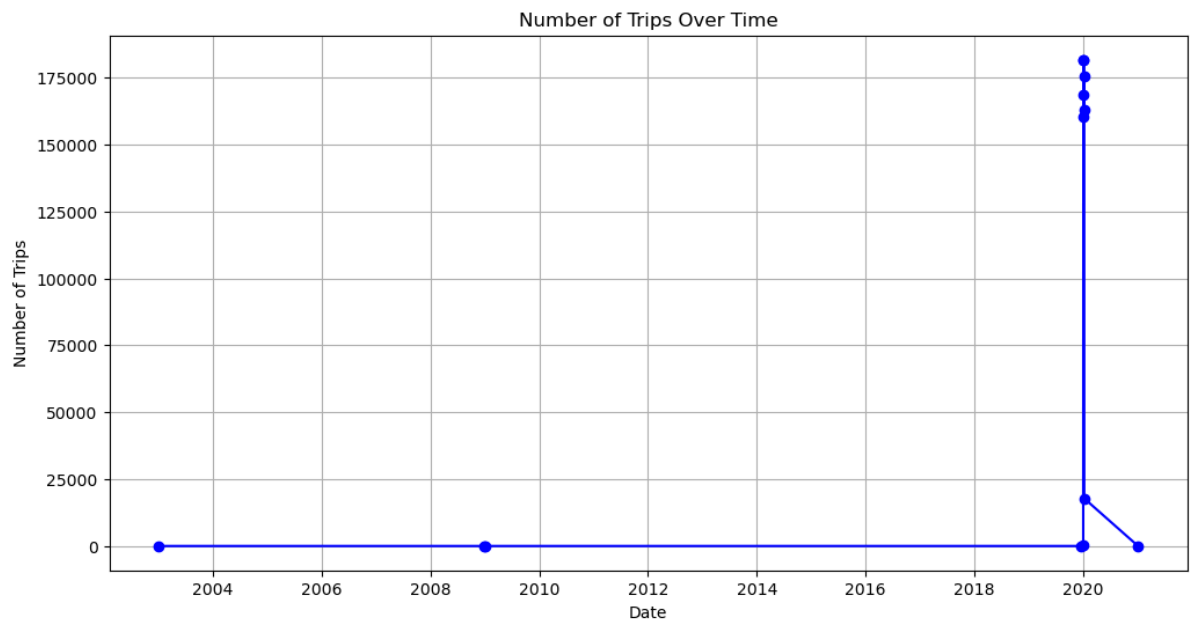
# Convert pickup datetime to datetime object with custom format
data['tpep_pickup_datetime'] = pd.to_datetime(data['tpep_pickup_datetime'], format='%Y-%m-%d %H:%M:%S')

# Extract date and time components
data['pickup_date'] = data['tpep_pickup_datetime'].dt.date
data['pickup_time'] = data['tpep_pickup_datetime'].dt.time

# Group by pickup date and count number of trips
trips_per_date = data.groupby('pickup_date').size()

# Plotting the number of trips over time
plt.figure(figsize=(12, 6))
trips_per_date.plot(kind='line', marker='o', color='b')
plt.title('Number of Trips Over Time')
plt.xlabel('Date')
plt.ylabel('Number of Trips')
plt.grid(True)
plt.show()

```



```
In [8]: import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset
data = pd.read_csv("C:/Users/91830/OneDrive/Desktop/yellow_tripdata_2020-01.csv")

# Convert pickup datetime to datetime object with custom format
data['tpep_pickup_datetime'] = pd.to_datetime(data['tpep_pickup_datetime'], format='%m/%d/%Y %H:%M:%S')

# Extract date and time components
data['pickup_date'] = data['tpep_pickup_datetime'].dt.date
data['pickup_time'] = data['tpep_pickup_datetime'].dt.time


# Group by pickup date and calculate number of trips and total trip distance
trips_distance_per_date = data.groupby('pickup_date').agg({'trip_distance': 'sum',
trips_distance_per_date.rename(columns={'passenger_count': 'num_trips'}, inplace=True)

# Plotting the number of trips and total trip distance over time
fig, ax1 = plt.subplots(figsize=(12, 6))

color = 'tab:blue'
ax1.set_xlabel('Date')
ax1.set_ylabel('Number of Trips', color=color)
ax1.plot(trips_distance_per_date.index, trips_distance_per_date['num_trips'], color=color)
ax1.tick_params(axis='y', labelcolor=color)

ax2 = ax1.twinx()
color = 'tab:red'
ax2.set_ylabel('Total Trip Distance', color=color)
ax2.plot(trips_distance_per_date.index, trips_distance_per_date['trip_distance'], color=color)
ax2.tick_params(axis='y', labelcolor=color)

plt.title('Number of Trips and Total Trip Distance Over Time')
plt.grid(True)
fig.tight_layout()
plt.show()
```



```
In [9]: import pandas as pd

# Load the dataset
data = pd.read_csv("C:/Users/91830/OneDrive/Desktop/yellow_tripdata_2020-01.csv")

# Convert pickup and drop-off datetime to datetime objects with custom format
data['tpep_pickup_datetime'] = pd.to_datetime(data['tpep_pickup_datetime'], format='%m/%d/%H%M')
data['tpep_dropoff_datetime'] = pd.to_datetime(data['tpep_dropoff_datetime'], format='%m/%d/%H%M')

# Extract date, time, and location components
data['pickup_date'] = data['tpep_pickup_datetime'].dt.date
data['pickup_time'] = data['tpep_pickup_datetime'].dt.time
data['dropoff_time'] = data['tpep_dropoff_datetime'].dt.time

# Group by pickup date and pickup/dropoff locations
repeated_rides = data.groupby(['pickup_date', 'PULocationID', 'DOLocationID']).filter(
    lambda x: x['PULocationID'] == x['DOLocationID']

# Display the repeated rides with pickup and drop-off locations, pickup time, and dropoff time
print(repeated_rides[['pickup_date', 'pickup_time', 'dropoff_time', 'PULocationID', 'DOLocationID']])
```

	pickup_date	pickup_time	dropoff_time	PULocationID	DOLocationID
0	2020-01-01	00:28:00	00:33:00	238	239
1	2020-01-01	00:35:00	00:43:00	239	238
2	2020-01-01	00:47:00	00:53:00	238	238
3	2020-01-01	00:55:00	01:00:00	238	151
4	2020-01-01	00:01:00	00:04:00	193	193
...	...	...	...	...	...
1048570	2020-01-07	07:33:00	07:38:00	229	140
1048571	2020-01-07	07:40:00	07:50:00	140	162
1048572	2020-01-07	07:19:00	07:28:00	48	161
1048573	2020-01-07	07:41:00	07:45:00	233	162
1048574	2020-01-07	07:39:00	07:43:00	79	137

[1027701 rows x 5 columns]

```
In [10]: import pandas as pd

# Load the dataset
data = pd.read_csv("C:/Users/91830/OneDrive/Desktop/yellow_tripdata_2020-01.csv")

# Group by pickup Location ID and count the occurrences
pickup_counts = data['PULocationID'].value_counts()

# Get the top 10 most frequent pickup Locations
top_pickup_locations = pickup_counts.head(10)
```

```
# Display the top pickup locations with their counts
print("Top 10 Most Frequent Pickup Locations:")
for location_id, count in top_pickup_locations.items():
    print(f"Location ID: {location_id}, Count: {count}")
```

Top 10 Most Frequent Pickup Locations:

```
Location ID: 132, Count: 47923
Location ID: 237, Count: 42010
Location ID: 161, Count: 40230
Location ID: 236, Count: 38638
Location ID: 186, Count: 38565
Location ID: 230, Count: 37981
Location ID: 48, Count: 35705
Location ID: 142, Count: 32580
Location ID: 162, Count: 31990
Location ID: 79, Count: 30128
```

In [11]: `import pandas as pd`

```
# Load the dataset
data = pd.read_csv("C:/Users/91830/OneDrive/Desktop/yellow_tripdata_2020-01.csv")

# Calculate profits for each ride
data['profits'] = data['total_amount'] - data['fare_amount']

# Sort the rides based on profits in descending order
top_profitable_rides = data.sort_values(by='profits', ascending=False)

# Display the top 10 profitable rides
print("Top 10 Profitable Rides:")
print(top_profitable_rides[['VendorID', 'tpep_pickup_datetime', 'tpep_dropoff_datetime', 'trip_distance', 'total_amount', 'profits']])
```

Top 10 Profitable Rides:

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	trip_distance	total_amount	profits
530908	1	04-01-2020 05:33	04-01-2020 05:34	0.40	965.80	913.80
55306	1	01-01-2020 04:06	01-01-2020 04:22	1.50	465.30	453.80
217700	2	02-01-2020 12:18	02-01-2020 12:36	4.19	352.30	336.80
1024290	2	06-01-2020 22:10	06-01-2020 22:10	0.00	1040.39	240.39
29954	2	01-01-2020 02:40	01-01-2020 02:57	4.02	247.89	231.89
1036479	1	07-01-2020 04:18	07-01-2020 05:37	44.00	374.96	196.46
586584	1	04-01-2020 14:55	04-01-2020 14:55	0.00	192.00	189.50
965333	1	06-01-2020 16:54	06-01-2020 20:00	86.00	581.06	182.56
811777	1	05-01-2020 17:09	05-01-2020 18:19	48.40	300.41	174.91
283321	1	02-01-2020 18:24	02-01-2020 18:32	6.40	165.00	165.00

	total_amount	profits
530908	965.80	913.80
55306	465.30	453.80
217700	352.30	336.80
1024290	1040.39	240.39
29954	247.89	231.89
1036479	374.96	196.46
586584	192.00	189.50
965333	581.06	182.56
811777	300.41	174.91
283321	165.00	165.00

In [17]: `import pandas as pd`

```
# Assuming your cleaned DataFrame is named 'cleaned_data'

# Specify the path where you want to save the CSV file
output_path = "C:/Users/91830/OneDrive/Desktop/clean_data.csv"
```

```
# Save the DataFrame as a CSV file  
data.to_csv(output_path, index=False)
```

In [ ]: